

Final Project Code

Homayoon Fotros

Libraries & Data Pre-processing

```
library(rio) ## importing data
library(tidyverse) ## data wrangling
library(corrplot) ## correlation plot
library(boot) ## bootstrap
library(gridExtra) ## grid visualization
library(glmnet) ## glm models
library(caret) ## various ML tools & models
library(ROCR) ## ROC-AUC score
library(leaps) ## model selection
library(bestglm) ## extended model selection
#library(GGally) ## optional: full features visualization
```

```
dt <- import('breastR2.xlsx') ## Make sure the data set file is in the working directory
head(dt[,1:9]) ## excluding the response variable
```

```
##   Age      BMI Glucose Insulin      HOMA  Leptin Adiponectin Resistin  MCP.1
## 1  48 23.50000      70   2.707 0.4674087  8.8071    9.702400   7.99585 417.114
## 2  83 20.69049      92   3.115 0.7068973  8.8438    5.429285   4.06405 468.786
## 3  82 23.12467      91   4.498 1.0096511 17.9393   22.432040   9.27715 554.697
## 4  68 21.36752      77   3.226 0.6127249  9.8827    7.169560  12.76600 928.220
## 5  86 21.11111      92   3.549 0.8053864  6.6994    4.819240  10.57635 773.920
## 6  49 22.85446      92   3.226 0.7320869  6.8317   13.679750  10.31760 530.410
```

```
## The positive response (diagnosed with BC) is coded as 1
dt$Classification <- ifelse(dt$Classification==1, 0,1)
table(dt$Classification)
```

```
##
##  0  1
## 52 64
```

```
str(dt) ## All the features are continuous numerical variables
```

```
## 'data.frame':   116 obs. of  10 variables:
##  $ Age      : num  48 83 82 68 86 49 89 76 73 75 ...
##  $ BMI      : num  23.5 20.7 23.1 21.4 21.1 ...
##  $ Glucose   : num  70 92 91 77 92 92 77 118 97 83 ...
##  $ Insulin   : num  2.71 3.12 4.5 3.23 3.55 ...
##  $ HOMA      : num  0.467 0.707 1.01 0.613 0.805 ...
##  $ Leptin    : num  8.81 8.84 17.94 9.88 6.7 ...
##  $ Adiponectin : num  9.7 5.43 22.43 7.17 4.82 ...
##  $ Resistin  : num  8 4.06 9.28 12.77 10.58 ...
##  $ MCP.1     : num  417 469 555 928 774 ...
```

```
## $ Classification: num 0 0 0 0 0 0 0 0 0 0 ...
```

```
summary(dt)
```

```
##      Age      BMI      Glucose      Insulin
## Min.   :24.0   Min.   :18.37   Min.    : 60.00   Min.    : 2.432
## 1st Qu.:45.0   1st Qu.:22.97   1st Qu.: 85.75   1st Qu.: 4.359
## Median :56.0   Median :27.66   Median : 92.00   Median : 5.925
## Mean   :57.3   Mean    :27.58   Mean    : 97.79   Mean    :10.012
## 3rd Qu.:71.0   3rd Qu.:31.24   3rd Qu.:102.00   3rd Qu.:11.189
## Max.   :89.0   Max.    :38.58   Max.    :201.00   Max.    :58.460
##      HOMA      Leptin      Adiponectin      Resistin
## Min.   : 0.4674   Min.    : 4.311   Min.    : 1.656   Min.    : 3.210
## 1st Qu.: 0.9180   1st Qu.:12.314   1st Qu.: 5.474   1st Qu.: 6.882
## Median : 1.3809   Median :20.271   Median : 8.353   Median :10.828
## Mean    : 2.6950   Mean     :26.615   Mean     :10.181   Mean     :14.726
## 3rd Qu.: 2.8578   3rd Qu.:37.378   3rd Qu.:11.816   3rd Qu.:17.755
## Max.    :25.0503   Max.     :90.280   Max.     :38.040   Max.     :82.100
##      MCP.1      Classification
## Min.    : 45.84   Min.     :0.0000
## 1st Qu.:269.98   1st Qu.:0.0000
## Median :471.32   Median :1.0000
## Mean     :534.65   Mean      :0.5517
## 3rd Qu.:700.09   3rd Qu.:1.0000
## Max.     :1698.44   Max.      :1.0000
```

```
set.seed(123) ## for replication
```

```
## Setting aside the validation holdout set (15% of the data)
```

```
train_test <- createDataPartition(dt$Classification, times = 1, p = 0.15, list = F)
```

```
holdout_set <- dt[train_test,]
```

```
dt <- dt[-train_test,]
```

```
## Number of instances in the training set:
```

```
nrow(dt)
```

```
## [1] 98
```

```
## The training set is fairly split between the two classes of response
```

```
table(dt$Classification)
```

```
##
```

```
## 0 1
```

```
## 44 54
```

```
table(holdout_set$Classification)
```


```
##
```

```
## 0 1
```

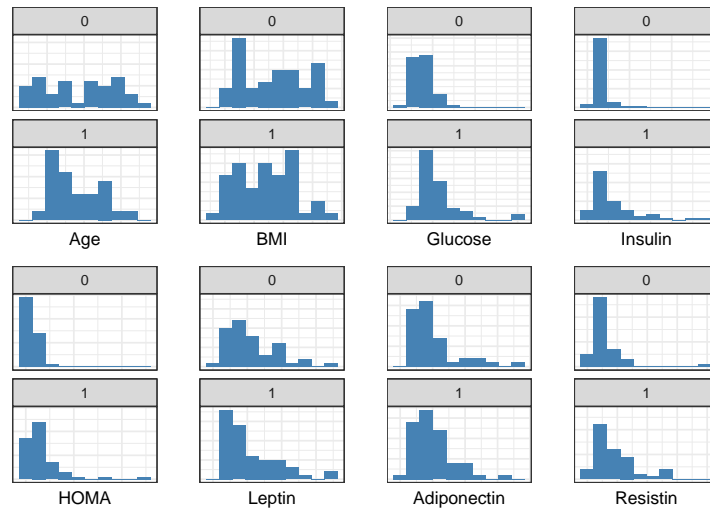
```
## 8 10
```

Exploratory Data Analysis (EDA)

Correlation between features and the response

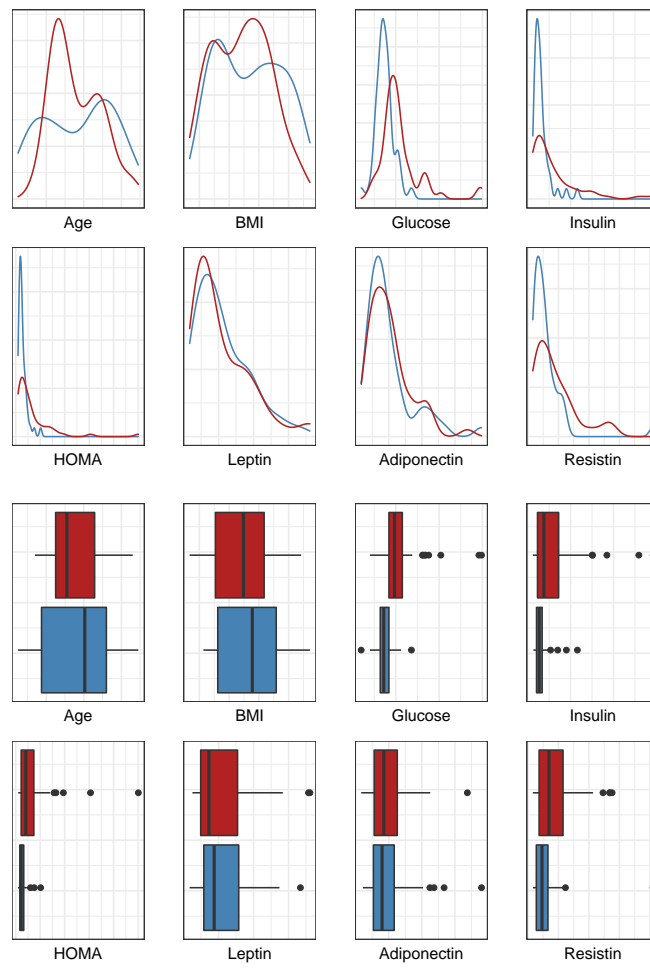


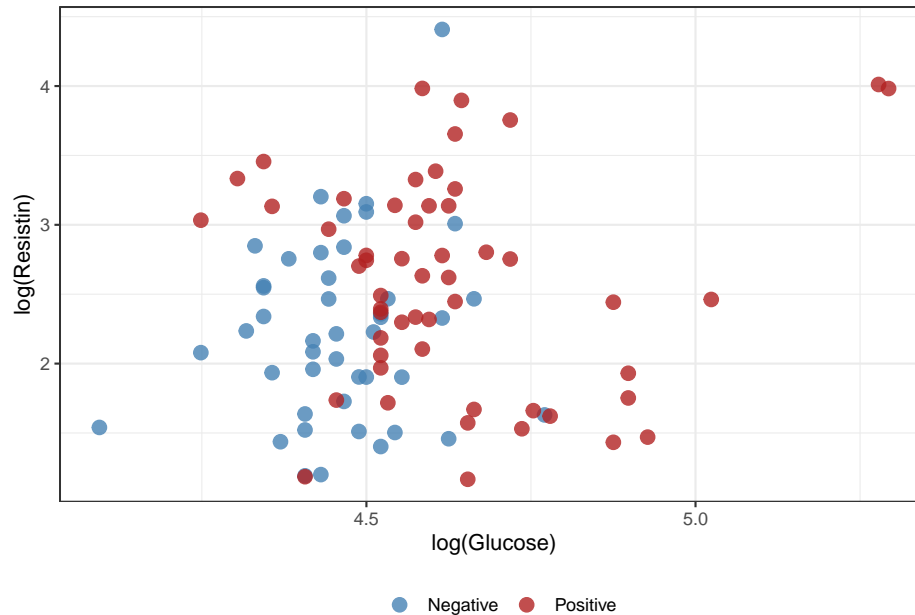
Distribution of Instances for each Feature by Class



Frequency & Range of Instances for each Feature by Class

Red: Positive (w/ BC) | Blue: Negative (wo/ BC)





Data Preparation

Log & Scale Transformation

```
dt_old <- dt ## backup of train data set before transformations
#dt <- dt_old

log_trans <- function(x) {sapply(x, log)}
scale_trans <- function(x){sapply(x, scale)}

dt[,3:9] <- log_trans(dt[,3:9])
dt[,1:9] <- data.frame(scale_trans(dt[,1:9]))

holdout_old <- holdout_set ## backup of test data set before transformations
#holdout_set<- holdout_old

holdout_set[,3:9] <- log_trans(holdout_set[,3:9])
holdout_set[,1:9] <- data.frame(scale_trans(holdout_set[,1:9]))
```

Setting Up Feature Matrix & Dataframe

```
# Train data matrix and data frame
mdl.mat <- model.matrix(Classification ~ . + I(Age^2) + I(BMI^2), data = dt)
train_dat <- data.frame(mdl.mat, Classification = factor(dt$Classification),
                        check.names = F)[,-1]

# Validation data matrix and data frame
mdl.tst <- model.matrix(Classification ~ . + I(Age^2) + I(BMI^2), data = holdout_set)
test_dat <- data.frame(mdl.tst, Classification = factor(holdout_set$Classification),
                      check.names = F)[,-1]
```

Classification Models

Linear Logistic Models

Logistic Model w/ & Wo/ Bootstrap

```
## Simple Logistic Model
mdl.log.simp <- glm(Classification ~ ., family = 'binomial', data = train_dat)
coef.simp <- coef(mdl.log.simp)
#summary(mdl.log.simp)
exp(coef(mdl.log.simp))

## (Intercept)      Age      BMI      Glucose      Insulin      HOMA
## 3.6636467 0.5248574 0.1995530 0.0000000 0.0000000      Inf
##      Leptin Adiponectin      Resistin      MCP.1 `I(Age^2)` `I(BMI^2)`
## 1.5915139 1.7569791 2.2827981 1.2534737 0.2128077 0.6976718

## Testing for models' fit (reduced (without Age^2 & BMI^2) vs. full)
mdl.log.red <- glm(Classification ~ ., family = 'binomial', data = train_dat[,c(-10,-11)])
#summary(mdl.log.red)
anova(mdl.log.red, mdl.log.simp, test = 'F')

## Analysis of Deviance Table
##
## Model 1: Classification ~ Age + BMI + Glucose + Insulin + HOMA + Leptin +
##      Adiponectin + Resistin + MCP.1
## Model 2: Classification ~ Age + BMI + Glucose + Insulin + HOMA + Leptin +
##      Adiponectin + Resistin + MCP.1 + `I(Age^2)` + `I(BMI^2)`
##      Resid. Df Resid. Dev Df Deviance      F      Pr(>F)
## 1          88      81.000
## 2          86      62.696 2    18.303 9.1517 0.000106 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

# Pseudo R-Squared values
DescTools::PseudoR2(mdl.log.simp, which = c('CoxSnell', 'McFadden'))

## CoxSnell  McFadden
## 0.5210244 0.5350134

## Bootstrap Logistic Model
## Setting up Bootstrap function
set.seed(123)
fn_est <- function(data, idx){
  mdl.boot <- glm(Classification ~ ., family = 'binomial', data = data[idx,])
  coefs <- mdl.boot$coefficients
  return(coefs)
}

L <- 1000 ## bootstrap samples = 1000
bts_est <- boot(data = train_dat, statistic = fn_est, R=L)

## storing bootstrap results in a matrix
boot.mat <- matrix(0, nrow = L, ncol=length(colnames(mdl.mat)))
for (i in 1:L) boot.mat[i,] <- bts_est$t[i,]

boot.mat <- apply(boot.mat, MARGIN = 2, median, na.rm = T) ## Median of bootstrapped coefficients
```

```

names(boot.mat) <- colnames(mdl.mat)

## setting up a function to calculate prob. score out of logit score
logit.pred <- function(data, coefs, thresh=0.5){
  factor(ifelse(1/(1+exp(-(data %*% coefs)))>=thresh, 1,0))
}

y_prd_simp <- logit.pred(mdl.mat, coef.simp) ## prediction of simple logistic model
y_prd_boot <- logit.pred(mdl.mat, boot.mat) ## prediction of logit bootstrap model

## setting up a function to calculate classification metrics
metrics <- function(y_true, y_hat){
  tr_accu <- mean(y_true == y_hat)
  tr_sens <- sensitivity(y_true, y_hat, positive = '1')
  tr_spc <- specificity(y_true, y_hat, negative = '0')
  tr_f1 <- F_meas(y_true, y_hat)
  prd <- prediction(as.numeric(y_hat), labels = y_true)
  tr_auc <- as.numeric(performance(prd, measure="auc")@y.values)
  data.frame(Accuracy = tr_accu, Recall=tr_sens, Specificity=tr_spc, F1=tr_f1, AUC=tr_auc)
}

## this data frame store our models' prediction metrics on the training set
model.perf <- data.frame(Model="Logistic",
  round(metrics(train_dat$Classification, y_prd_simp),2))
model.perf <- rbind(model.perf, data.frame(Model="Logistic (Boot)",
  round(metrics(train_dat$Classification, y_prd_boot),2)))

```

Lasso CV Models

```

## Lasso L00V
set.seed(123)
mdl.loigtL00V <- cv.glmnet(x = mdl.mat[,-1], y = train_dat$Classification,
  family=binomial, nfolds = nrow(train_dat))

## Coefficients (L00V Lasso)
Lasso_LV_Coefs <- predict(mdl.loigtL00V, newx = mdl.mat[,-1], type='coefficient')
coef_lass_LV <- data.frame(Var = c('Intercept',
  Lasso_LV_Coefs@Dimnames[[1]][-1])[c(Lasso_LV_Coefs@i+1)],
  `Lasso L00V` = round(Lasso_LV_Coefs@x, 3))

coef_lass_LV_ac <- rbind(coef_lass_LV,
  data.frame(Var = colnames(mdl.mat)[-1][c(!((colnames(mdl.mat)[-1])%in% coef_la

## Performance on the training set
y_prd_lgL00V <- factor(ifelse(predict(mdl.loigtL00V, newx = mdl.mat[,-1],
  type='response')>0.5, 1,0))

## Lasso K-fold(10)
set.seed(123)
mdl.loigtKfold <- cv.glmnet(x = mdl.mat[,-1], y = train_dat$Classification,
  family=binomial, nfolds = 10)

#Coefficients (k-fold Lasso) [uncomment to run]

```

```

Lasso_K10_Coefs <- predict(mdl.loigtKfold, newx = mdl.mat[,-1],type='coefficient')

coef_lass_k10 <- data.frame(Var = c('Intercept',
                                   Lasso_K10_Coefs@Dimnames[[1]][-1])[c(Lasso_K10_Coefs@i+1)],
  `Lasso K-fold` = round(Lasso_K10_Coefs@x, 3))

coef_lass_k10_ac <- rbind(coef_lass_k10,
  data.frame(Var = colnames(mdl.mat)[-1][c(!((colnames(mdl.mat)[-1])%in% coef_1

## Performance on the training set
y_prd_lgKfold <- factor(ifelse(predict(mdl.loigtKfold, newx = mdl.mat[,-1],
                                     type='response')>0.5, 1,0))

## coefficients of the Lasso models:
coefs_lasso_final <- merge(coef_lass_k10_ac, coef_lass_LV_ac, by='Var')

## storing the models' performance in the metrics data frame
model.perf <- rbind(model.perf, data.frame(
  Model='Logistic (L00V)',round(metrics(train_dat$Classification,
                                     y_prd_lgL00V),2)))
model.perf <- rbind(model.perf, data.frame(
  Model='Logistic (Kfold=10)', round(metrics(train_dat$Classification,
                                     y_prd_lgKfold),2)))

```

Logit Subset Models

```

## Subset Criteria (BestGLM)
set.seed(123)
lgs_subs <- function(data, ctn='AIC', meth='exhaustive'){
  mdl.bglm <- bestglm(data, family = binomial, IC=ctn, method = meth)
  tmp_names <- names(mdl.bglm$BestModel$coefficients)
  tmp_names <- case_when(
    tmp_names == "`I(Age^2)`" ~ "I(Age^2)",
    tmp_names == "`I(BMI^2)`" ~ "I(BMI^2)",
    TRUE ~ as.character(tmp_names)
  )
  coefs <- mdl.bglm$BestModel$coefficients
  names(coefs) <- tmp_names
  return(coefs)
}

coefs_ex_AIC <- lgs_subs(data=train_dat)
coefs_ex_BIC <- lgs_subs(data=train_dat,ctn='BIC')
coefs_fw_AIC <- lgs_subs(data=train_dat, meth = 'forward')
coefs_fw_BIC <- lgs_subs(data=train_dat,ctn = 'BIC', meth = 'forward')
coefs_bw_AIC <- lgs_subs(data=train_dat, meth = 'backward')
coefs_bw_BIC <- lgs_subs(data=train_dat,ctn = 'BIC', meth = 'backward')

y_prd_ex_AIC <- logit.pred(mdl.mat[,names(coefs_ex_AIC)],
  coefs_ex_AIC)
y_prd_ex_BIC <- logit.pred(mdl.mat[,names(coefs_ex_BIC)],
  coefs_ex_BIC)
y_prd_fw_AIC <- logit.pred(mdl.mat[,names(coefs_fw_AIC)],
  coefs_fw_AIC)

```



```

y_prd_fw_BIC <- logit.pred(mdl.mat[,names(coefs_fw_BIC)],
                           coefs_fw_BIC)
y_prd_bw_AIC <- logit.pred(mdl.mat[,names(coefs_bw_AIC)],
                           coefs_bw_AIC)
y_prd_bw_BIC <- logit.pred(mdl.mat[,names(coefs_bw_BIC)],
                           coefs_bw_BIC)
## Metrics of All Subset Models
cbind(Model=c('Best subset - AIC', 'Best subset - BIC', 'Forward - AIC',
              'Forward - BIC', 'Backward - AIC', 'Backward - BIC' ),
      round(rbind(metrics(train_dat$Classification, y_prd_ex_AIC ),
                    metrics(train_dat$Classification, y_prd_ex_BIC ),
                    metrics(train_dat$Classification, y_prd_fw_AIC ),
                    metrics(train_dat$Classification, y_prd_fw_BIC ),
                    metrics(train_dat$Classification, y_prd_bw_AIC ),
                    metrics(train_dat$Classification, y_prd_bw_BIC )),2))

##           Model Accuracy Recall Specificity   F1   AUC
## 1 Best subset - AIC      0.87   0.87         0.86 0.85 0.86
## 2 Best subset - BIC      0.84   0.84         0.83 0.81 0.83
## 3 Forward - AIC         0.87   0.87         0.86 0.85 0.86
## 4 Forward - BIC         0.84   0.84         0.83 0.81 0.83
## 5 Backward - AIC        0.87   0.87         0.86 0.85 0.86
## 6 Backward - BIC        0.84   0.84         0.83 0.81 0.83
## Coefficients of all selected logistic models
coefs_all_logit <- merge(data.frame(Var = c('Intercept', colnames(mdl.mat)[-1]),
                                       Simple.Logit = round(exp(coef.simp),3),
                                       row.names = NULL, check.names = F),
                        merge (coefs_lasso_final,
                              data.frame(Var = c('Intercept', names(coefs_fw_AIC)[-1]),
                                           Subset = round(exp(coefs_fw_AIC),3), row.names = NULL),
                              by='Var', all.x = T, all.y = T, ),
                        all.x = T, all.y = T, by='Var')

coefs_all_logit[is.na(coefs_all_logit)] <-0
coefs_all_logit[rank(c('Intercept', colnames(mdl.mat)[-1])),] %>%
  data.frame(row.names = NULL)

##           Var Simple.Logit Lasso.K.fold Lasso.L00V Subset
## 1 Intercept      3.664         0.734         0.786 2.059
## 2 Age            0.525         0.000         0.000 0.000
## 3 BMI            0.200        -0.245        -0.293 0.267
## 4 Glucose        0.000         0.739         0.783 0.000
## 5 Insulin        0.000         0.000         0.000 0.000
## 6 HOMA           Inf         0.149         0.176   Inf
## 7 Leptin         1.592         0.000         0.000 0.000
## 8 Adiponectin    1.757         0.000         0.000 1.810
## 9 Resistin       2.283         0.218         0.256 2.670
## 10 MCP.1         1.253         0.000         0.000 0.000
## 11 I(Age^2)      0.213        -0.490        -0.526 0.236
## 12 I(BMI^2)      0.698         0.000        -0.010 0.000

model.perf <- rbind(model.perf, data.frame(
  Model='Logistic (Fwd Sel.)', round(metrics(train_dat$Classification, y_prd_fw_AIC ),2)))

```

Non-Parametric Models

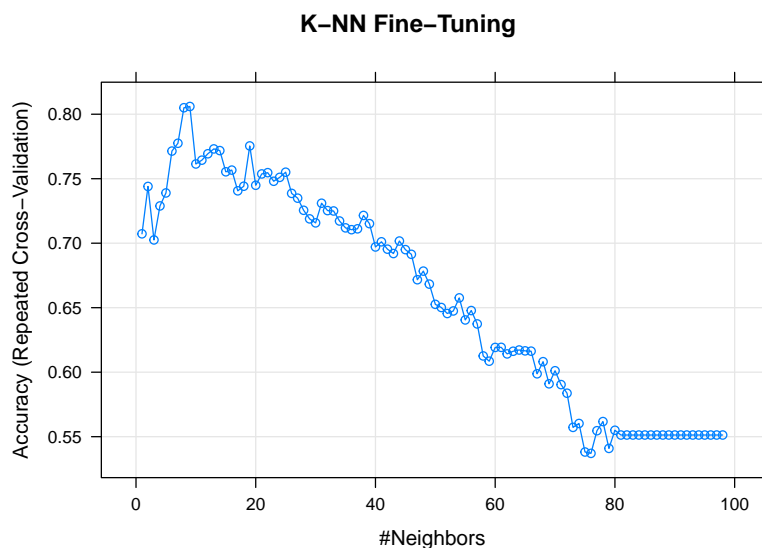
k-NN

```
set.seed(123)
## set CV (10 fold, 3 repeats)
tr <- trainControl(method = 'repeatedcv', number = 10, repeats = 3)

k_grid <- data.frame(k=seq(1,nrow(dt),1)) # tune for number of neighbors: 1 to n
clf_knn <- train(Classification ~. , data = train_dat,
                 trControl = tr, method = 'knn', tuneGrid = k_grid)
# Best k (number of neighbors) based on cross-validated accuracy
clf_knn$bestTune

## k
## 9 9

plot(clf_knn, main = 'K-NN Fine-Tuning')
```



```
y_prd_knn <- predict(clf_knn, newdata = train_dat, type='raw')
model.perf <- rbind(model.perf, data.frame(
  Model='K-NN', round(metrics(train_dat$Classification, y_prd_knn),2)))
```

Random Forest

```
rf_res <- data.frame(N_Trees = NULL, mtry = NULL, Accuracy =NULL)

for (nt in seq(500, 3000, by = 500)){

  set.seed(123)
  tg <- expand.grid(.mtry = seq(1,11,by=1))
  clf_rf <- train(Classification ~. , data = train_dat,
                 trControl = tr, method = 'rf', tuneGrid = tg, ntree=nt)

  rf_res <- rbind(rf_res, data.frame(N_Trees = clf_rf$finalModel$ntree,
                                    mtry = clf_rf$finalModel$mtry,
```

```

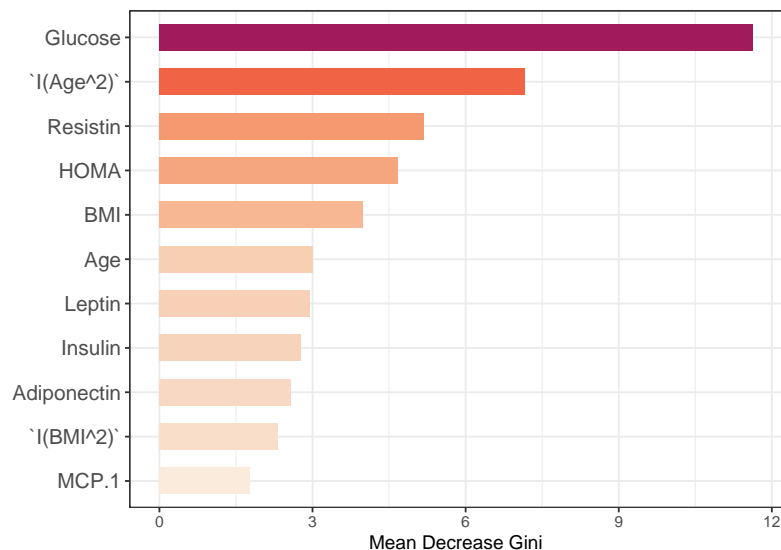
Accuracy = max(clf_rf$results$Accuracy)))
}
## result of grid search for RF models
rf_res

##   N_Trees mtry Accuracy
## 1     500   10 0.7394949
## 2    1000    6 0.7398653
## 3    1500    6 0.7361616
## 4    2000    6 0.7324579
## 5    2500    6 0.7394949
## 6    3000    6 0.7321549

## so we use mtry = 6 and ntree=1000

clf_rf <- train(Classification ~. , data = train_dat, method = 'rf',
                 tuneGrid = expand.grid(.mtry = 6), ntree=1000)
rf_feats <- clf_rf$finalModel$importance[,1]
ggplot() +
  geom_col(aes(y=reorder(names(rf_feats), rf_feats), x= rf_feats, fill=rf_feats), width = 0.6) +
  scale_fill_viridis(option = 'rocket', direction = -1, begin=0.4, end=1) +
  labs(x='Mean Decrease Gini', y='') +
  theme(legend.position = 'none', axis.text.y = element_text(size=12))

```



```

y_pred_rf <- predict(clf_rf, newdata = train_dat, type = 'raw')

model.perf <- rbind(model.perf, data.frame(
  Model='RF', round(metrics(train_dat$Classification, clf_rf$finalModel$predicted),2)))

```

SVM (Linear & Radial Kernel)

```

### SVM Linear
set.seed(123)

clf_svm_l <- train(Classification~., data = train_dat,

```

```

trControl = tr, method = "svmLinear", preProcess = c("center","scale"),
tuneGrid = expand.grid(C=seq(0.1,10, by=0.5)))

y_pred_svm_l <- predict(clf_svm_l, newdata = train_dat, type = 'raw')
clf_svm_l$bestTune

```

```

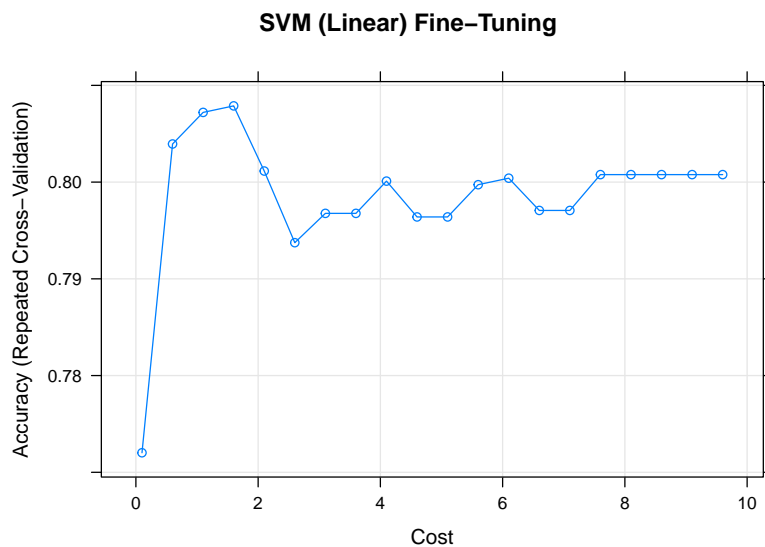
##      C
## 4 1.6

```

```

plot(clf_svm_l, main='SVM (Linear) Fine-Tuning')

```



```

model.perf <- rbind(model.perf, data.frame(
  Model='SVM (Linear)', round(metrics(train_dat$Classification, y_pred_svm_l),2)))

```

```

### SVM Radial

```

```

set.seed(123)

```

```

clf_svm_r <- train(Classification~., data = train_dat,
  trControl = tr, method = "svmRadial",
  preProcess = c("center","scale"),
  tuneGrid = expand.grid(C=seq(1,100, length.out = 10),
    sigma= seq(0.01, 0.05, length.out = 10)))
clf_svm_r$bestTune

```

```

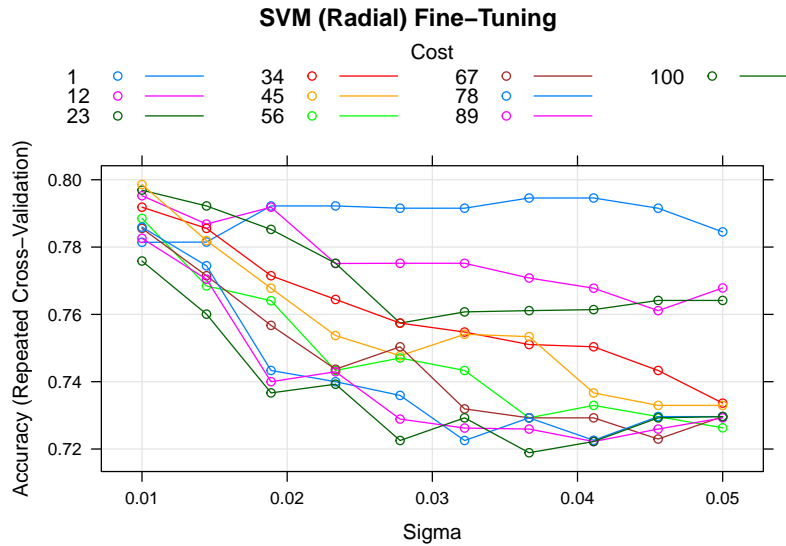
##      sigma C
## 41 0.01 45

```

```

#clf_svm_r$finalModel
plot(clf_svm_r, main='SVM (Radial) Fine-Tuning')

```



```
y_pred_svm_r <- predict(clf_svm_r, newdata = train_dat, type = 'raw')

model.perf <- rbind(model.perf, data.frame(
  Model='SVM (Radial)', round(metrics(train_dat$Classification, y_pred_svm_r),2)))
```

```
model.perf
```

##	Model	Accuracy	Recall	Specificity	F1	AUC
## 1	Logistic	0.87	0.87	0.86	0.85	0.86
## 2	Logistic (Boot)	0.64	0.73	0.58	0.65	0.65
## 3	Logistic (LOOV)	0.83	0.80	0.86	0.79	0.82
## 4	Logistic (Kfold=10)	0.82	0.79	0.86	0.78	0.81
## 5	Logistic (Fwd Sel.)	0.87	0.87	0.86	0.85	0.86
## 6	K-NN	0.86	0.87	0.84	0.84	0.86
## 7	RF	0.76	0.75	0.76	0.71	0.75
## 8	SVM (Linear)	0.88	0.86	0.90	0.86	0.87
## 9	SVM (Radial)	0.94	0.93	0.95	0.93	0.94

Applying on the Validation/Holdout Set

```
## Logistic Simple Model
test_lm <- logit.pred(mdl.tst,coef.simp)
tst.perf <- data.frame(Model="Logistic (Base)",round(metrics(test_dat$Classification, test_lm),2))

## Logistic Bootstrap
test_boot <- logit.pred(mdl.tst,boot.mat)
tst.perf <- rbind(tst.perf, data.frame(Model="Logistic (Boot)",
  round(metrics(test_dat$Classification, test_boot),2)))

## Lasso
test_lgL00V <- factor(ifelse(predict(mdl.loigtL00V, newx = mdl.tst[, -1],
  type='response')>0.5, 1,0))

test_lgKfold <- factor(ifelse(predict(mdl.loigtKfold, newx = mdl.tst[, -1],
```

```

                                type='response')>0.5, 1,0))

tst.perf <- rbind(tst.perf, data.frame(
  Model='Lasso (LOOV)', round(metrics(test_dat$Classification, test_lgLOOV),2)))
tst.perf <- rbind(tst.perf, data.frame(
  Model='Lasso (K-fold)', round(metrics(test_dat$Classification, test_lgKfold),2)))

## FW. Sel.
test_fw_AIC <- logit.pred(mdl.tst[,names(coefs_fw_AIC)],
                        coefs_fw_AIC)
tst.perf <- rbind(tst.perf, data.frame(
  Model='Logistic (Fwd Sel.)', round(metrics(test_dat$Classification, test_fw_AIC ),2)))

## K-NN
test_knn <- predict(clf_knn, newdata = test_dat, type='raw')
tst.perf <- rbind(tst.perf, data.frame(
  Model='K-NN', round(metrics(test_dat$Classification, test_knn),2)))

## RF
test_rf <- predict(clf_rf, newdata = test_dat, type = 'raw')
tst.perf <- rbind(tst.perf, data.frame(
  Model='RF', round(metrics(test_dat$Classification, test_rf),2)))

## SVM Linear
test_svm_l <- predict(clf_svm_l, newdata = test_dat, type = 'raw')
tst.perf <- rbind(tst.perf, data.frame(
  Model='SVM (Linear)', round(metrics(test_dat$Classification, test_svm_l),2)))

## SVM Radial
test_svm_r <- predict(clf_svm_r, newdata = test_dat, type = 'raw')
tst.perf <- rbind(tst.perf, data.frame(
  Model='SVM (Radial)', round(metrics(test_dat$Classification, test_svm_r),2)))

tst.perf

```

##	Model	Accuracy	Recall	Specificity	F1	AUC
## 1	Logistic (Base)	0.50	0.55	0.43	0.40	0.49
## 2	Logistic (Boot)	0.44	0.50	0.33	0.29	0.42
## 3	Lasso (LOOV)	0.67	0.67	0.67	0.57	0.65
## 4	Lasso (K-fold)	0.67	0.67	0.67	0.57	0.65
## 5	Logistic (Fwd Sel.)	0.50	0.55	0.43	0.40	0.49
## 6	K-NN	0.83	0.82	0.86	0.80	0.82
## 7	RF	0.72	0.86	0.64	0.74	0.74
## 8	SVM (Linear)	0.72	0.73	0.71	0.67	0.71
## 9	SVM (Radial)	0.89	1.00	0.80	0.89	0.90